

## Redes neuronales artificiales con Keras

*Este notebook contiene todo el código fuente requerido para la solución de los talleres propuestos.*



Ejecutar en Google Colab ([https://colab.research.google.com/github/ageron/handson-ml2/blob/master/10\\_neural\\_nets\\_with\\_keras.ipynb](https://colab.research.google.com/github/ageron/handson-ml2/blob/master/10_neural_nets_with_keras.ipynb))

## Configuración

Primero, importemos algunos módulos comunes, asegurémonos de que Matplotlib traza las figuras en línea y preparemos una función para guardar las figuras. También verifiquemos que Python 3.5 o posterior esté instalado (aunque Python 2.x puede funcionar, está obsoleto, por lo que recomendamos encarecidamente que use Python 3 en su lugar), así como Scikit-Learn  $\geq 0.20$  y TensorFlow  $\geq 2.0$ .

In [116]:

```

1  # Python ≥3.5 es requerido
2  import sys
3  assert sys.version_info >= (3, 5)
4
5  # Scikit-Learn ≥0.20 es requerido
6  import sklearn
7  assert sklearn.__version__ >= "0.20"
8
9  try:
10     # %tensorflow_version solo existe en Colab.
11     %tensorflow_version 2.x
12 except Exception:
13     pass
14
15 # TensorFlow ≥2.0 es requerido
16 import tensorflow as tf
17 assert tf.__version__ >= "2.0"
18
19 # Importar Librerías comunes
20 import numpy as np
21 import os
22
23 # para que la salida de este notebook sea estable en todas las ejecuciones
24 np.random.seed(42)
25
26 # Para dibujar figuras estéticas
27 %matplotlib inline
28 import matplotlib as mpl
29 import matplotlib.pyplot as plt
30 mpl.rc('axes', labelsize=14)
31 mpl.rc('xtick', labelsize=12)
32 mpl.rc('ytick', labelsize=12)
33
34 # En donde almacenar las figuras
35 PROJECT_ROOT_DIR = "."
36 CHAPTER_ID = "ann"
37 IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
38 os.makedirs(IMAGES_PATH, exist_ok=True)
39
40 def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300)
41     path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
42     print("Saving figure", fig_id)
43     if tight_layout:
44         plt.tight_layout()
45     plt.savefig(path, format=fig_extension, dpi=resolution)
46
47 # Ignorar Las advertencias inútiles (consulte el número 5998 de SciPy)
48 import warnings
49 warnings.filterwarnings(action="ignore", message="^internal gelsd")

```

## Perceptrones

**Nota:** establecemos `max_iter` y `tol` explícitamente para evitar advertencias sobre el hecho de que su valor predeterminado cambiará en futuras versiones de Scikit-Learn.

```
In [117]: 1 import numpy as np
          2 from sklearn.datasets import load_iris
          3 from sklearn.linear_model import Perceptron
          4
          5 iris = load_iris()
          6 X = iris.data[:, (2, 3)] # Largo del pétalo, ancho del pétalo
          7 y = (iris.target == 0).astype(np.int)
          8
          9 per_clf = Perceptron(max_iter=1000, tol=1e-3, random_state=42)
         10 per_clf.fit(X, y)
         11
         12 y_pred = per_clf.predict([[2, 0.5]])
```

```
In [118]: 1 y_pred
```


```
Out[118]: array([1])
```

```

In [119]: 1 a = -per_clf.coef_[0][0] / per_clf.coef_[0][1]
2 b = -per_clf.intercept_ / per_clf.coef_[0][1]
3
4 axes = [0, 5, 0, 2]
5
6 x0, x1 = np.meshgrid(
7     np.linspace(axes[0], axes[1], 500).reshape(-1, 1),
8     np.linspace(axes[2], axes[3], 200).reshape(-1, 1),
9 )
10 X_new = np.c_[x0.ravel(), x1.ravel()]
11 y_predict = per_clf.predict(X_new)
12 zz = y_predict.reshape(x0.shape)
13
14 plt.figure(figsize=(10, 4))
15 plt.plot(X[y==0, 0], X[y==0, 1], "bs", label="No es Iris-Setosa")
16 plt.plot(X[y==1, 0], X[y==1, 1], "yo", label="Iris-Setosa")
17
18 plt.plot([axes[0], axes[1]], [a * axes[0] + b, a * axes[1] + b], "k-", linewidth=2)
19 from matplotlib.colors import ListedColormap
20 custom_cmap = ListedColormap(['#9898ff', '#fafab0'])
21
22 plt.contourf(x0, x1, zz, cmap=custom_cmap)
23 plt.xlabel("Longitud Pétalo", fontsize=14)
24 plt.ylabel("Ancho Pétalo", fontsize=14)
25 plt.legend(loc="lower right", fontsize=14)
26 plt.axis(axes)
27
28 save_fig("perceptron_iris_plot")
29 plt.show()

```

Saving figure perceptron\_iris\_plot



```

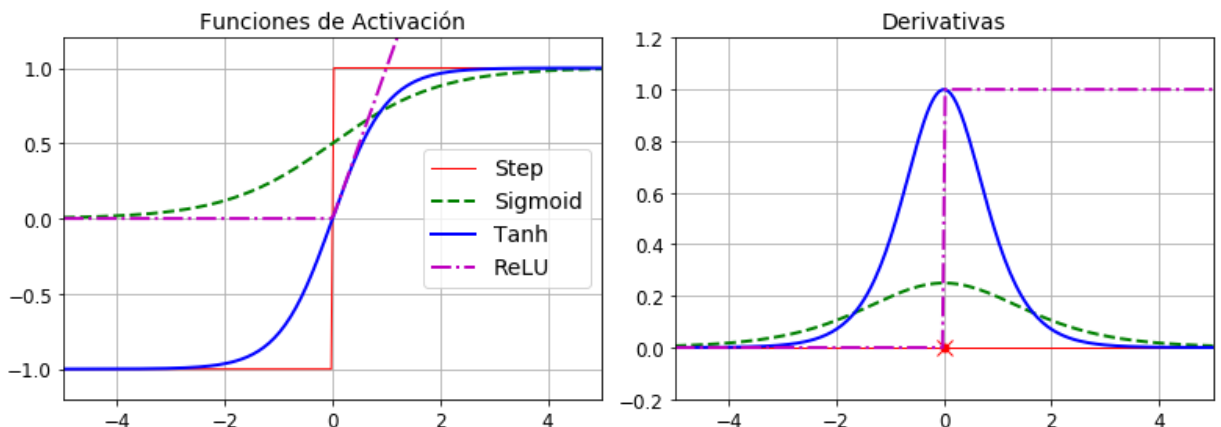
In [120]: 1 # Funciones de Activación

```

```
In [121]: 1 def sigmoid(z):
2         return 1 / (1 + np.exp(-z))
3
4 def relu(z):
5         return np.maximum(0, z)
6
7 def derivative(f, z, eps=0.000001):
8         return (f(z + eps) - f(z - eps))/(2 * eps)
```

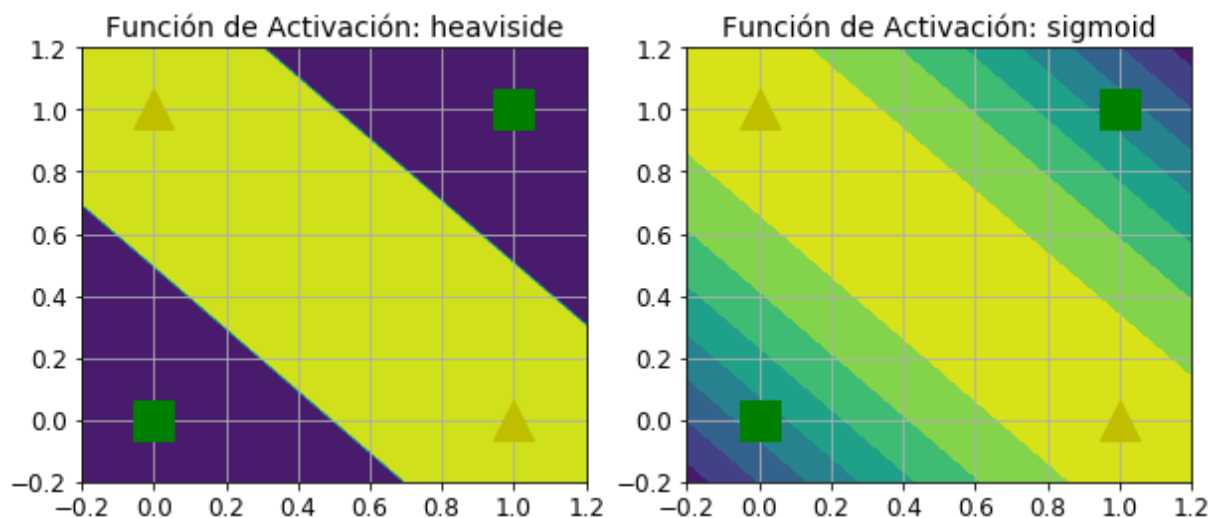
```
In [122]: 1 z = np.linspace(-5, 5, 200)
2
3 plt.figure(figsize=(11,4))
4
5 plt.subplot(121)
6 plt.plot(z, np.sign(z), "r-", linewidth=1, label="Step")
7 plt.plot(z, sigmoid(z), "g--", linewidth=2, label="Sigmoid")
8 plt.plot(z, np.tanh(z), "b-", linewidth=2, label="Tanh")
9 plt.plot(z, relu(z), "m-.", linewidth=2, label="ReLU")
10 plt.grid(True)
11 plt.legend(loc="center right", fontsize=14)
12 plt.title("Funciones de Activación", fontsize=14)
13 plt.axis([-5, 5, -1.2, 1.2])
14
15 plt.subplot(122)
16 plt.plot(z, derivative(np.sign, z), "r-", linewidth=1, label="Step")
17 plt.plot(0, 0, "ro", markersize=5)
18 plt.plot(0, 0, "rx", markersize=10)
19 plt.plot(z, derivative(sigmoid, z), "g--", linewidth=2, label="Sigmoid")
20 plt.plot(z, derivative(np.tanh, z), "b-", linewidth=2, label="Tanh")
21 plt.plot(z, derivative(relu, z), "m-.", linewidth=2, label="ReLU")
22 plt.grid(True)
23 #plt.legend(loc="center right", fontsize=14)
24 plt.title("Derivativas", fontsize=14)
25 plt.axis([-5, 5, -0.2, 1.2])
26
27 save_fig("activation_functions_plot")
28 plt.show()
```

Saving figure activation\_functions\_plot



```
In [123]: 1 def heaviside(z):
2         return (z >= 0).astype(z.dtype)
3
4 def mlp_xor(x1, x2, activation=heaviside):
5         return activation(-activation(x1 + x2 - 1.5) + activation(x1 + x2 - 0.5))
```

```
In [124]: 1 x1s = np.linspace(-0.2, 1.2, 100)
2 x2s = np.linspace(-0.2, 1.2, 100)
3 x1, x2 = np.meshgrid(x1s, x2s)
4
5 z1 = mlp_xor(x1, x2, activation=heaviside)
6 z2 = mlp_xor(x1, x2, activation=sigmoid)
7
8 plt.figure(figsize=(10,4))
9
10 plt.subplot(121)
11 plt.contourf(x1, x2, z1)
12 plt.plot([0, 1], [0, 1], "gs", markersize=20)
13 plt.plot([0, 1], [1, 0], "y^", markersize=20)
14 plt.title("Función de Activación: heaviside", fontsize=14)
15 plt.grid(True)
16
17 plt.subplot(122)
18 plt.contourf(x1, x2, z2)
19 plt.plot([0, 1], [0, 1], "gs", markersize=20)
20 plt.plot([0, 1], [1, 0], "y^", markersize=20)
21 plt.title("Función de Activación: sigmoid", fontsize=14)
22 plt.grid(True)
```



## Construyendo un clasificador de Imágenes

Primero, importemos TensorFlow y Keras.

```
In [125]: 1 import tensorflow as tf
2 from tensorflow import keras
```

```
In [126]: 1 tf.__version__
```

```
Out[126]: '2.3.1'
```

```
In [127]: 1 keras.__version__
```

```
Out[127]: '2.4.0'
```

Comencemos cargando el conjunto de datos de moda MNIST. Keras tiene una serie de funciones para cargar conjuntos de datos populares en `keras.datasets`. El conjunto de datos ya está dividido entre un conjunto de entrenamiento y un conjunto de prueba, pero puede ser útil dividir el conjunto de entrenamiento más para tener un conjunto de validación:

```
In [128]: 1 fashion_mnist = keras.datasets.fashion_mnist  
2 (X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()
```

El conjunto de entrenamiento contiene 60.000 imágenes en escala de grises, cada una de 28x28 píxeles:

```
In [129]: 1 X_train_full.shape
```

```
Out[129]: (60000, 28, 28)
```

Cada intensidad de píxel se representa como un byte (de 0 a 255):

```
In [130]: 1 X_train_full.dtype
```


```
Out[130]: dtype('uint8')
```

Dividamos el conjunto de entrenamiento completo en un conjunto de validación y un conjunto de entrenamiento (más pequeño). También escalamos las intensidades de píxeles hasta el rango 0-1 y las convertimos en flotantes, dividiéndolas por 255.

```
In [131]: 1 X_valid, X_train = X_train_full[:5000] / 255., X_train_full[5000:] / 255.  
2 y_valid, y_train = y_train_full[:5000], y_train_full[5000:]  
3 X_test = X_test / 255.
```

Puede trazar una imagen usando la función `imshow()` de Matplotlib, con un 'binario' mapa de color:

```
In [132]: 1 plt.imshow(X_train[0], cmap="binary")
          2 plt.axis('off')
          3 plt.show()
```



Las etiquetas son los ID de clase (representados como uint8), de 0 a 9:

```
In [133]: 1 y_train
```

```
Out[133]: array([4, 0, 7, ..., 3, 0, 5], dtype=uint8)
```

Aquí están los nombres de clase correspondientes:

```
In [134]: 1 class_names = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat",
          2                  "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"]
```

Entonces, la primera imagen del conjunto de entrenamiento es un abrigo:

```
In [135]: 1 class_names[y_train[0]]
```

```
Out[135]: 'Coat'
```

El conjunto de validación contiene 5000 imágenes y el conjunto de prueba contiene 10000 imágenes:

```
In [136]: 1 X_valid.shape
```

```
Out[136]: (5000, 28, 28)
```



In [137]: 1 X\_test.shape

Out[137]: (10000, 28, 28)

Echemos un vistazo a una muestra de las imágenes en el conjunto de datos:

```
In [138]: 1 n_rows = 4
2 n_cols = 10
3 plt.figure(figsize=(n_cols * 1.2, n_rows * 1.2))
4 for row in range(n_rows):
5     for col in range(n_cols):
6         index = n_cols * row + col
7         plt.subplot(n_rows, n_cols, index + 1)
8         plt.imshow(X_train[index], cmap="binary", interpolation="nearest")
9         plt.axis('off')
10        plt.title(class_names[y_train[index]], fontsize=12)
11 plt.subplots_adjust(wspace=0.2, hspace=0.5)
12 save_fig('fashion_mnist_plot', tight_layout=False)
13 plt.show()
```

Saving figure fashion\_mnist\_plot



```
In [139]: 1 model = keras.models.Sequential()
2 model.add(keras.layers.Flatten(input_shape=[28, 28]))
3 model.add(keras.layers.Dense(300, activation="relu"))
4 model.add(keras.layers.Dense(100, activation="relu"))
5 model.add(keras.layers.Dense(10, activation="softmax"))
```

```
In [140]: 1 keras.backend.clear_session()
2 np.random.seed(42)
3 tf.random.set_seed(42)
```

```
In [141]: 1 model = keras.models.Sequential([
2         keras.layers.Flatten(input_shape=[28, 28]),
3         keras.layers.Dense(300, activation="relu"),
4         keras.layers.Dense(100, activation="relu"),
5         keras.layers.Dense(10, activation="softmax")
6     ])
```

```
In [142]: 1 model.layers
```

```
Out[142]: [<tensorflow.python.keras.layers.core.Flatten at 0x299f2ae9448>,
<tensorflow.python.keras.layers.core.Dense at 0x299f2ae95c8>,
<tensorflow.python.keras.layers.core.Dense at 0x299f2ae9ac8>,
<tensorflow.python.keras.layers.core.Dense at 0x299f2ad4188>]
```

```
In [143]: 1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 784)	0
-----		
dense (Dense)	(None, 300)	235500
-----		
dense_1 (Dense)	(None, 100)	30100
-----		
dense_2 (Dense)	(None, 10)	1010
=====		
Total params: 266,610		
Trainable params: 266,610		
Non-trainable params: 0		

```
In [144]: 1 keras.utils.plot_model(model, "my_fashion_mnist_model.png", show_shapes=True
('Failed to import pydot. You must `pip install pydot` and install graphviz (ht
tps://graphviz.gitlab.io/download/), ', 'for `pydotprint` to work.')
```

```
In [145]: 1 hidden1 = model.layers[1]
2         hidden1.name
```

```
Out[145]: 'dense'
```

```
In [146]: 1 model.get_layer(hidden1.name) is hidden1
```

```
Out[146]: True
```

```
In [147]: 1 weights, biases = hidden1.get_weights()
```

```
In [148]: 1 weights
```

```
Out[148]: array([[ 0.02448617, -0.00877795, -0.02189048, ..., -0.02766046,  
                  0.03859074, -0.06889391],  
                 [ 0.00476504, -0.03105379, -0.0586676 , ...,  0.00602964,  
                  -0.02763776, -0.04165364],  
                 [-0.06189284, -0.06901957,  0.07102345, ..., -0.04238207,  
                  0.07121518, -0.07331658],  
                 ...,  
                 [-0.03048757,  0.02155137, -0.05400612, ..., -0.00113463,  
                  0.00228987,  0.05581069],  
                 [ 0.07061854, -0.06960931,  0.07038955, ..., -0.00384101,  
                  0.00034875,  0.02878492],  
                 [-0.06022581,  0.01577859, -0.02585464, ..., -0.00527829,  
                  0.00272203, -0.06793761]], dtype=float32)
```

```
In [149]: 1 weights.shape
```

Out[149]: (784, 300)

```
In [150]: 1 biases
```

[illegible]

```
In [151]: 1 biases.shape
```

Out[151]: (300,)

```
In [152]: 1 model.compile(loss="sparse_categorical_crossentropy",
2               optimizer="sgd",
3               metrics=["accuracy"])
```

Esto es equivalente a:

```
model.compile(loss=keras.losses.sparse_categorical_crossentropy,
```

```
optimizer=keras.optimizers.SGD(),  
metrics=[keras.metrics.sparse_categorical_accuracy])
```

In [153]:

```
1 history = model.fit(X_train, y_train, epochs=30,
2                     validation_data=(X_valid, y_valid))
```

```
Epoch 1/30
1719/1719 [=====] - 6s 4ms/step - loss: 0.7237 - accur
acy: 0.7644 - val_loss: 0.5207 - val_accuracy: 0.8234
Epoch 2/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.4843 - accur
acy: 0.8318 - val_loss: 0.4345 - val_accuracy: 0.8538
Epoch 3/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.4393 - accur
acy: 0.8454 - val_loss: 0.5341 - val_accuracy: 0.7988
Epoch 4/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.4126 - accur
acy: 0.8565 - val_loss: 0.3915 - val_accuracy: 0.8644
Epoch 5/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3940 - accur
acy: 0.8618 - val_loss: 0.3748 - val_accuracy: 0.8690
Epoch 6/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3753 - accur
acy: 0.8677 - val_loss: 0.3707 - val_accuracy: 0.8728
Epoch 7/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3633 - accur
acy: 0.8715 - val_loss: 0.3623 - val_accuracy: 0.8720
Epoch 8/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3519 - accur
acy: 0.8750 - val_loss: 0.3848 - val_accuracy: 0.8624
Epoch 9/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3416 - accur
acy: 0.8792 - val_loss: 0.3588 - val_accuracy: 0.8704
Epoch 10/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3324 - accur
acy: 0.8819 - val_loss: 0.3427 - val_accuracy: 0.8780
Epoch 11/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3243 - accur
acy: 0.8835 - val_loss: 0.3433 - val_accuracy: 0.8786
Epoch 12/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3151 - accur
acy: 0.8866 - val_loss: 0.3310 - val_accuracy: 0.8820
Epoch 13/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3083 - accur
acy: 0.8885 - val_loss: 0.3262 - val_accuracy: 0.8888
Epoch 14/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3024 - accur
acy: 0.8914 - val_loss: 0.3387 - val_accuracy: 0.8774
Epoch 15/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2950 - accur
acy: 0.8939 - val_loss: 0.3205 - val_accuracy: 0.8864
Epoch 16/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2892 - accur
acy: 0.8972 - val_loss: 0.3083 - val_accuracy: 0.8908
Epoch 17/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2841 - accur
acy: 0.8977 - val_loss: 0.3546 - val_accuracy: 0.8740
Epoch 18/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2780 - accur
acy: 0.9000 - val_loss: 0.3138 - val_accuracy: 0.8902
```

```
Epoch 19/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2729 - accur
acy: 0.9021 - val_loss: 0.3130 - val_accuracy: 0.8898
Epoch 20/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2678 - accur
acy: 0.9035 - val_loss: 0.3271 - val_accuracy: 0.8804
Epoch 21/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2626 - accur
acy: 0.9056 - val_loss: 0.3069 - val_accuracy: 0.8918
Epoch 22/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2578 - accur
acy: 0.9072 - val_loss: 0.2971 - val_accuracy: 0.8960
Epoch 23/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2538 - accur
acy: 0.9082 - val_loss: 0.2986 - val_accuracy: 0.8936
Epoch 24/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2485 - accur
acy: 0.9105 - val_loss: 0.3073 - val_accuracy: 0.8882
Epoch 25/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2447 - accur
acy: 0.9121 - val_loss: 0.2970 - val_accuracy: 0.8954
Epoch 26/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2409 - accur
acy: 0.9136 - val_loss: 0.3055 - val_accuracy: 0.8888
Epoch 27/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2366 - accur
acy: 0.9155 - val_loss: 0.3019 - val_accuracy: 0.8952
Epoch 28/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2331 - accur
acy: 0.9167 - val_loss: 0.2993 - val_accuracy: 0.8930
Epoch 29/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2287 - accur
acy: 0.9182 - val_loss: 0.3052 - val_accuracy: 0.8892
Epoch 30/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2255 - accur
acy: 0.9187 - val_loss: 0.3032 - val_accuracy: 0.8930
```

```
In [154]: 1 history.params
```

```
Out[154]: {'verbose': 1, 'epochs': 30, 'steps': 1719}
```

```
In [155]: 1 print(history.epoch)
```


```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21,
22, 23, 24, 25, 26, 27, 28, 29]
```

```
In [156]: 1 history.history.keys()
```

```
Out[156]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [157]: 1 import pandas as pd
2
3 pd.DataFrame(history.history).plot(figsize=(8, 5))
4 plt.grid(True)
5 plt.gca().set_ylim(0, 1)
6 save_fig("keras_learning_curves_plot")
7 plt.show()
```

Saving figure keras\_learning\_curves\_plot



```
In [158]: 1 model.evaluate(X_test, y_test)
```

313/313 [=====] - 1s 2ms/step - loss: 0.3377 - accuracy: 0.8827

```
Out[158]: [0.33773481845855713, 0.8827000260353088]
```

```
In [159]: 1 X_new = X_test[:3]
          2 y_proba = model.predict(X_new)
          3 y_proba.round(2)
```

WARNING:tensorflow:5 out of the last 7 calls to <function Model.make\_predict\_function.<locals>.predict\_function at 0x00000299F2547AF8> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has experimental\_relax\_shapes=True option that relaxes argument shapes that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/tutorials/customization/performance#python\\_or\\_tensor\\_args](https://www.tensorflow.org/tutorials/customization/performance#python_or_tensor_args) ([https://www.tensorflow.org/tutorials/customization/performance#python\\_or\\_tensor\\_args](https://www.tensorflow.org/tutorials/customization/performance#python_or_tensor_args)) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) ([https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function)) for more details.

```
Out[159]: array([[0. , 0. , 0. , 0. , 0. , 0.01, 0. , 0.02, 0. , 0.96],
                 [0. , 0. , 0.99, 0. , 0.01, 0. , 0. , 0. , 0. , 0. ],
                 [0. , 1. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ]],
          dtype=float32)
```

```
In [160]: 1 y_pred = model.predict_classes(X_new)
          2 y_pred
```

```
Out[160]: array([9, 2, 1], dtype=int64)
```

```
In [161]: 1 np.array(class_names)[y_pred]
```

```
Out[161]: array(['Ankle boot', 'Pullover', 'Trouser'], dtype='<U11')
```


```
In [162]: 1 y_new = y_test[:3]
          2 y_new
```

```
Out[162]: array([9, 2, 1], dtype=uint8)
```



```
In [163]: 1 plt.figure(figsize=(7.2, 2.4))
2 for index, image in enumerate(X_new):
3     plt.subplot(1, 3, index + 1)
4     plt.imshow(image, cmap="binary", interpolation="nearest")
5     plt.axis('off')
6     plt.title(class_names[y_test[index]], fontsize=12)
7 plt.subplots_adjust(wspace=0.2, hspace=0.5)
8 save_fig('fashion_mnist_images_plot', tight_layout=False)
9 plt.show()
```

Saving figure fashion\_mnist\_images\_plot



## Regresión MLP

Carguemos, dividamos y escalemos el conjunto de datos de viviendas de California:

```
In [164]: 1 from sklearn.datasets import fetch_california_housing
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler
4
5 housing = fetch_california_housing()
6
7 X_train_full, X_test, y_train_full, y_test = train_test_split(housing.data,
8 X_train, X_valid, y_train, y_valid = train_test_split(X_train_full, y_train_
9
10 scaler = StandardScaler()
11 X_train = scaler.fit_transform(X_train)
12 X_valid = scaler.transform(X_valid)
13 X_test = scaler.transform(X_test)
```

```
In [165]: 1 np.random.seed(42)
2 tf.random.set_seed(42)
```

```
In [166]: 1 model = keras.models.Sequential([
2         keras.layers.Dense(30, activation="relu", input_shape=X_train.shape[1:])
3         keras.layers.Dense(1)
4     ])
5 model.compile(loss="mean_squared_error", optimizer=keras.optimizers.SGD(lr=1
6 history = model.fit(X_train, y_train, epochs=20, validation_data=(X_valid, y
7 mse_test = model.evaluate(X_test, y_test)
8 X_new = X_test[:3]
9 y_pred = model.predict(X_new)
```

```
Epoch 1/20
363/363 [=====] - 1s 3ms/step - loss: 1.6419 - val_loss: 0.8560
Epoch 2/20
363/363 [=====] - 1s 2ms/step - loss: 0.7047 - val_loss: 0.6531
Epoch 3/20
363/363 [=====] - 1s 2ms/step - loss: 0.6345 - val_loss: 0.6099
Epoch 4/20
363/363 [=====] - 1s 2ms/step - loss: 0.5977 - val_loss: 0.5658
Epoch 5/20
363/363 [=====] - 1s 2ms/step - loss: 0.5706 - val_loss: 0.5355
Epoch 6/20
363/363 [=====] - 1s 2ms/step - loss: 0.5472 - val_loss: 0.5173
Epoch 7/20
363/363 [=====] - 1s 2ms/step - loss: 0.5288 - val_loss: 0.5081
Epoch 8/20
363/363 [=====] - 1s 2ms/step - loss: 0.5130 - val_loss: 0.4799
Epoch 9/20
363/363 [=====] - 1s 2ms/step - loss: 0.4992 - val_loss: 0.4690
Epoch 10/20
363/363 [=====] - 1s 2ms/step - loss: 0.4875 - val_loss: 0.4656
Epoch 11/20
363/363 [=====] - 1s 2ms/step - loss: 0.4777 - val_loss: 0.4482
Epoch 12/20
363/363 [=====] - 1s 2ms/step - loss: 0.4688 - val_loss: 0.4479
Epoch 13/20
363/363 [=====] - ETA: 0s - loss: 0.463 - 1s 2ms/step - loss: 0.4615 - val_loss: 0.4296
Epoch 14/20
363/363 [=====] - 1s 2ms/step - loss: 0.4547 - val_loss: 0.4233
Epoch 15/20
363/363 [=====] - 1s 2ms/step - loss: 0.4488 - val_loss: 0.4176
Epoch 16/20
363/363 [=====] - 1s 2ms/step - loss: 0.4435 - val_loss: 0.4176
```

```

s: 0.4123
Epoch 17/20
363/363 [=====] - 1s 2ms/step - loss: 0.4389 - val_loss: 0.4071
Epoch 18/20
363/363 [=====] - 1s 2ms/step - loss: 0.4347 - val_loss: 0.4037
Epoch 19/20
363/363 [=====] - 1s 2ms/step - loss: 0.4306 - val_loss: 0.4000
Epoch 20/20
363/363 [=====] - 1s 2ms/step - loss: 0.4273 - val_loss: 0.3969
162/162 [=====] - 0s 1ms/step - loss: 0.4212
WARNING:tensorflow:6 out of the last 9 calls to <function Model.make_predict_function.<locals>.predict_function at 0x00000299F2540AF8> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has experimental_relax_shapes=True option that relaxes argument shapes that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/tutorials/customization/performance#python\_or\_tensor\_args (https://www.tensorflow.org/tutorials/customization/performance#python\_or\_tensor\_args) and https://www.tensorflow.org/api\_docs/python/tf/function (https://www.tensorflow.org/api\_docs/python/tf/function) for more details.


```

In [167]:

```

1 plt.plot(pd.DataFrame(history.history))
2 plt.grid(True)
3 plt.gca().set_ylim(0, 1)
4 plt.show()

```



In [168]:

```
1 y_pred
```

```

Out[168]: array([[0.3885664],
                  [1.6792021],
                  [3.1022797]], dtype=float32)

```

```
In [169]: 1 class PrintValTrainRatioCallback(keras.callbacks.Callback):
2         def on_epoch_end(self, epoch, logs):
3             print("\nval/train: {:.2f}".format(logs["val_loss"] / logs["loss"]))
```

```
In [170]: 1 val_train_ratio_cb = PrintValTrainRatioCallback()
2         history = model.fit(X_train, y_train, epochs=1,
3                             validation_data=(X_valid, y_valid),
4                             callbacks=[val_train_ratio_cb])
```

1/363 [.....] - ETA: 0s - loss: 0.8053WARNING:tensorflow:Callbacks method `on\_train\_batch\_begin` is slow compared to the batch time (batch time: 0.0000s vs `on\_train\_batch\_begin` time: 0.0010s). Check your callbacks.

WARNING:tensorflow:Callbacks method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0000s vs `on\_train\_batch\_end` time: 0.0014s). Check your callbacks.

354/363 [=====>.] - ETA: 0s - loss: 0.4212WARNING:tensorflow:Callbacks method `on\_test\_batch\_begin` is slow compared to the batch time (batch time: 0.0013s vs `on\_test\_batch\_begin` time: 0.0029s). Check your callbacks.

val/train: 0.93

363/363 [=====] - 1s 2ms/step - loss: 0.4240 - val\_loss: 0.3932

```
In [171]: 1 model.evaluate(X_test, y_test)
```

162/162 [=====] - 0s 2ms/step - loss: 0.4184

Out[171]: 0.4183996021747589